

If I Could Churn Back Time

... we wouldn't need these models, gnome sayin'?

Overview

For my capstone project, I have chosen to continue in a similar vein to my Phase 3 project with entirely different data. Churn, which is what it is called when a customer stops doing business with a company, is one of the largest issues facing almost every industry today. In context of subscription industries, for example, "churn" is a massive ongoing issue, as it is infinitely cheaper to keep existing customers than to lure in new ones. Therefore, establishing a procedure to keep customers from "churning" is of the utmost importance. The first step in this process is identifying the reasons behind customer churn, such as dissatisfaction, area of service provided, quality of service, and other such considerations. Following that, we will build a model that will accurately predict whether or not a customer will churn, and thereby provide the limitlessly valuable opportunity for a company to prevent that churn.

Business Understanding

As previously stated, it is much cheaper to keep existing customers than to continually lure new customers in. In other words, the longer a customer stays with a company, the more money that company stands to make from them. It seems fairly simple in that context, right? So it also follows that if there was a way to accurately predict whether or not a customer would churn, that method would be likely to save a company a boatload of money. And no, "boatload" in this case is not an exaggeration - the data we are using today comes from a fictional telecommunications company, but the very real company known as Verizon walked away from 2023 with \$134 billion in total operating revenue. So our goal here was quite simple - build a model that could accurately predict when a customer would stop doing business with a company, and provide that company an opportunity to stop that customer from churning, thereby saving them money.

Data Understanding & Preparation

The data we've used for this endeavor comes to us from IBM, and is a publicly available dataset created for this exact purpose - to help deal with churn problems, and allow students like me to cut our teeth on such before we enter the wider world of data science. This sample data module tracks a fictional Telco company's customer churn based on a variety of possible factors, such as gender, monthly charges, and usage information, as well as whether the customer churned or not. Due to the nature of the data we used, any location based features were excluded because the data was all centered in California specifically.

In [1]: |

Multiple

Lines

Phone

Service

24

72

11

Yes

Yes

Yes

Yes

Yes

No

Yes

Yes

No

phone

service

Out[3]:

```
TIIIPUI L PAITUAS AS PU
import numpy as np
from sklearn.model_selection import train_test_split, cross_val_score, Randomi
from sklearn.pipeline import Pipeline as ImbPipeline
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.dummy import DummyClassifier
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
from sklearn.metrics import classification_report, ConfusionMatrixDisplay, Rod
from sklearn.metrics import make_scorer
from sklearn.metrics import accuracy_score, recall_score, precision_score, f1_
from sklearn import tree
from sklearn.preprocessing import OrdinalEncoder
from sklearn.preprocessing import LabelEncoder
from sklearn.ensemble import RandomForestClassifier
```

Our first step in data understanding is to load in our various Excel spreadsheets and transform them into Pandas dataframes to be explored. Our first exploratory steps will inform the rest of the process from cleaning the data to modeling, all the way to our final recommendations.

```
In [2]: churn_df = pd.read_excel('Data/CustomerChurn.xlsx')
In [3]: churn_df
```

Partner Dependents Tenure

Senior

Citizen

Customer

6840-

RESVB 2234-

XADUH

4801-

JZAZL

No

No

No

Yes

Yes

Yes

LoyaltyID

0	318537	7590- VHVEG	No	Yes	No	1	No	No phone service
1	152148	5575- GNVDE	No	No	No	34	Yes	No
2	326527	3668- QPYBK	No	No	No	2	Yes	No
3	845894	7795- CFOCW	No	No	No	45	No	No phone service
4	503388	9237- HQITU	No	No	No	2	Yes	No

810338

230811

155157

7038

7039

7040

```
8361-
         7041
                 731782
                                       Yes
                                                Yes
                                                            No
                                                                      4
                                                                             Yes
                                                                                      Yes
                            LTMKD
                             3186-
         7042
                 353947
                                       No
                                                No
                                                            No
                                                                     66
                                                                             Yes
                                                                                      No
                             AJIEK
        7043 rows × 21 columns
In [4]:
         churn df.info()
       <class 'pandas.core.frame.DataFrame'>
       RangeIndex: 7043 entries, 0 to 7042
       Data columns (total 21 columns):
            Column
                                Non-Null Count
                                                 Dtype
        0
            LoyaltyID
                                7043 non-null
                                                 int64
        1
            Customer ID
                                7043 non-null
                                                 object
        2
            Senior Citizen
                                7043 non-null
                                                 object
        3
            Partner
                                7043 non-null
                                                 object
        4
            Dependents
                                7043 non-null
                                                object
        5
            Tenure
                                7043 non-null
                                                 int64
        6
            Phone Service
                                7043 non-null
                                                object
        7
            Multiple Lines
                                7043 non-null
                                                 object
        8
                                                 object
            Internet Service
                                7043 non-null
        9
            Online Security
                                7043 non-null
                                                 object
        10 Online Backup
                                7043 non-null
                                                 object
        11 Device Protection 7043 non-null
                                                 object
        12 Tech Support
                                7043 non-null
                                                 object
        13 Streaming TV
                                7043 non-null
                                                 object
                                7043 non-null
        14 Streaming Movies
                                                 object
        15 Contract
                                7043 non-null
                                                 object
        16 Paperless Billing
                                7043 non-null
                                                 object
            Payment Method
                                7043 non-null
        17
                                                 object
        18
            Monthly Charges
                                7043 non-null
                                                 float64
            Total Charges
        19
                                7043 non-null
                                                object
                                                 object
        20 Churn
                                7043 non-null
       dtypes: float64(1), int64(2), object(18)
       memory usage: 1.1+ MB
In [5]:
         churn_df.isna().sum()
Out[5]:
         LoyaltyID
                               0
         Customer ID
                               0
         Senior Citizen
                               0
         Partner
                               0
         Dependents
                               0
                               0
         Tenure
         Phone Service
                               0
                               0
         Multiple Lines
                               0
         Internet Service
                               0
         Online Security
         Online Backup
                               0
         Device Protection
                               0
         Tech Support
                               0
                               0
         Streaming TV
         Streaming Movies
                               0
                               0
         Contract
         Paperless Billing
                               0
         Payment Method
                               0
         Monthly Charges
                               0
                               0
         Total Charges
         Churn
```

dtype: int64

Right off the bat with our first dataframe, we can see that there are no null values and that our dataframe contains floats and ints (which are numbers) and objects. Something to note is that "Total Charges" is listed as an object and it likely ought to be an int or a float. After some fiddling, it became clear that this was going to be troublesome, but I found myself in luck - our third dataset has properly formatted information about charges for various services as well as total charges overall, so we will not need to keep that information in this dataset at all.

Out[7]:		Customer ID	Senior Citizen	Partner	Dependents	Tenure	Phone Service	Multiple Lines	Internet Service	Sı
	0	7590- VHVEG	No	Yes	No	1	No	No phone service	DSL	
	1	5575- GNVDE	No	No	No	34	Yes	No	DSL	
	2	3668- QPYBK	No	No	No	2	Yes	No	DSL	
	3	7795- CFOCW	No	No	No	45	No	No phone service	DSL	
	4	9237- HQITU	No	No	No	2	Yes	No	Fiber optic	
	•••									
	7038	6840- RESVB	No	Yes	Yes	24	Yes	Yes	DSL	
	7039	2234- XADUH	No	Yes	Yes	72	Yes	Yes	Fiber optic	
	7040	4801- JZAZL	No	Yes	Yes	11	No	No phone service	DSL	
	7041	8361- LTMKD	Yes	Yes	No	4	Yes	Yes	Fiber optic	
	7042	3186- AJIEK	No	No	No	66	Yes	No	Fiber optic	

7043 rows × 18 columns

Now, we will repeat these same steps with the other relevant datasets.

```
In [8]:
            telco churn = pd.read excel('Data/Telco customer churn.xlsx')
 In [9]:
           telco_churn
 Out[9]:
                                                                         Zip
                                                                 City
                  CustomerID Count Country
                                                     State
                                                                                 Lat Long
                                                                                             Latitude
                                                                       Code
                                         United
                                                                                33.964131,
                                                                 Los
                                    1
                                                                       90003
              0 3668-QPYBK
                                                 California
                                                                                            33.964131
                                                                              -118.272783
                                          States
                                                             Angeles
                                         United
                                                                 Los
                                                                                34.059281,
                                    1
                                                                       90005
                                                                                            34.059281
                  9237-HQITU
                                                 California
                                          States
                                                             Angeles
                                                                                -118.30742
                                         United
                                                                                34.048013,
                                                                 Los
                 9305-CDSKC
                                    1
                                                  California
                                                                       90006
                                                                                            34.048013
                                          States
                                                                              -118.293953
                                                             Angeles
                        7892-
                                         United
                                                                                34.062125,
                                                                 Los
              3
                                    1
                                                  California
                                                                       90010
                                                                                            34.062125
                       POOKP
                                          States
                                                                              -118.315709
                                                             Angeles
                                                                                34.039224,
                                         United
                                                                       90015
                                                 California
                                                                                            34.039224
                  0280-XJGEX
                                    1
                                          States
                                                             Angeles
                                                                              -118.266293
              •••
                        2569-
                                         United
                                                                                34.341737,
           7038
                                                                                            34.341737
                                    1
                                                  California
                                                             Landers
                                                                       92285
                      WGERO
                                          States
                                                                              -116.539416
                                         United
                                                                                34.667815,
           7039
                  6840-RESVB
                                                  California
                                                                       92301
                                                                                            34.667815
                                                            Adelanto
                                          States
                                                                              -117.536183
                        2234-
                                         United
                                                                                34.559882,
           7040
                                                  California
                                                              Amboy
                                                                       92304
                                                                                            34.559882
                      XADUH
                                          States
                                                                              -115.637164
                                         United
                                                             Angelus
                                                                                  34.1678,
           7041
                   4801-JZAZL
                                    1
                                                  California
                                                                       92305
                                                                                            34.167800
                                                                                -116.86433
                                          States
                                                                Oaks
                                         United
                                                                                34.424926,
                                                               Apple
                   3186-AJIEK
           7042
                                    1
                                                  California
                                                                       92308
                                                                                            34.424926
                                          States
                                                               Valley
                                                                              -117.184503
          7043 rows × 33 columns
In [10]:
           telco_churn.info()
         <class 'pandas.core.frame.DataFrame'>
         RangeIndex: 7043 entries, 0 to 7042
         Data columns (total 33 columns):
          #
               Column
                                    Non-Null Count
                                                       Dtype
          0
               CustomerID
                                    7043 non-null
                                                       object
          1
               Count
                                    7043 non-null
                                                       int64
          2
               Country
                                    7043 non-null
                                                      object
          3
               State
                                    7043 non-null
                                                       object
          4
               City
                                    7043 non-null
                                                       object
               Zip Code
                                    7043 non-null
                                                       int64
```

```
7043 non-null
                                      obiect
6
    Lat Long
7
                      7043 non-null
                                      float64
    Latitude
                                      float64
8
   Longitude
                      7043 non-null
9
    Gender
                      7043 non-null
                                      object
10 Senior Citizen
                      7043 non-null
                                      object
11 Partner
                      7043 non-null
                                      object
                                      object
12 Dependents
                      7043 non-null
13 Tenure Months
                                      int64
                      7043 non-null
14 Phone Service
                                      object
                      7043 non-null
15 Multiple Lines
                                     object
                      7043 non-null
16 Internet Service
                      7043 non-null
                                      object
17 Online Security
                      7043 non-null
                                      object
18 Online Backup
                      7043 non-null
                                     object
19 Device Protection 7043 non-null
                                      object
                      7043 non-null
20 Tech Support
                                      object
                      7043 non-null
                                      object
21 Streaming TV
22 Streaming Movies 7043 non-null
                                      object
23 Contract
                      7043 non-null
                                      object
24 Paperless Billing 7043 non-null
                                      object
25 Payment Method
                      7043 non-null
                                      object
26 Monthly Charges
                      7043 non-null
                                      float64
27 Total Charges
                      7043 non-null
                                      object
                      7043 non-null
28 Churn Label
                                      object
29 Churn Value
                      7043 non-null
                                      int64
30 Churn Score
                      7043 non-null
                                      int64
 31 CLTV
                      7043 non-null
                                      int64
 32 Churn Reason
                      1869 non-null
                                     object
dtypes: float64(3), int64(6), object(24)
```

memory usage: 1.8+ MB

There are a few immediate issues with this dataframe. Firstly, while "Churn Reason" seems like it might be very useful for modeling purposes, the column is also full of nulls, and so must be removed. Any location information can be removed because this dataset is based entirely in California, so any location-based reasonings would not be sound. We must also compare this dataframe to the first one and remove any duplicate or unnecessary columns, for ease of merging the two dataframes into one useful one. Many of these columns also "overlap" with our first dataframe, and therefore are unnecessary.

```
In [11]:
          telco_churn = telco_churn.drop(['Count', 'Country', 'State', 'City', 'Zip Code
In [12]:
          telco churn
```

Out[12]:		CustomerID	Gender	Churn Value	Churn Score	CLTV
	0	3668-QPYBK	Male	1	86	3239
	1	9237-HQITU	Female	1	67	2701
	2	9305-CDSKC	Female	1	86	5372
	3	7892-POOKP	Female	1	84	5003
	4	0280-XJGEX	Male	1	89	5340
	•••					
	7038	2569-WGERO	Female	0	45	5306
	7039	6840-RESVB	Male	0	59	2140
	7040	2234-XADUH	Female	0	71	5560

```
7041 4801-JZAZL Female 0 59 2793
7042 3186-AJIEK Male 0 38 5097
```

7043 rows × 5 columns

```
In [13]: telco_churn = telco_churn.rename(columns={'CustomerID': 'Customer ID'})
In [14]: telco_churn
```

\cap	+-	Γ1	4	
υu	L	LΤ	+	

	Customer ID	Gender	Churn Value	Churn Score	CLTV
0	3668-QPYBK	Male	1	86	3239
1	9237-HQITU	Female	1	67	2701
2	9305-CDSKC	Female	1	86	5372
3	7892-POOKP	Female	1	84	5003
4	0280-XJGEX	Male	1	89	5340
•••					
7038	2569-WGERO	Female	0	45	5306
7039	6840-RESVB	Male	0	59	2140
7040	2234-XADUH	Female	0	71	5560
7041	4801-JZAZL	Female	0	59	2793
7042	3186-AJIEK	Male	0	38	5097

7043 rows × 5 columns

Now, both of these dataframes are ready to be merged together, using Customer ID as a joining point since they are shared between them. This will add a few more features to our eventual modeling process, and I am specifically interested in gender as a feature here simply because it will be fascinating to see whether gender has any bearing on whether or not a customer stopped doing business with a company. I doubt it, but it is a unique viewpoint to consider.

We do have one final dataset that contains a few more services that could be used as features for modeling to determine whether those services (or lack thereof) have any impact on whether a customer will churn or not, so now, we will go through the same process one last time to assemble a third cleaned dataframe for our modeling purposes.

```
Tn [16]
```

```
services_df = pd.read_excel('Data/Telco_customer_churn_services.xlsx')
In [17]:
           services df
Out[17]:
                                                                       Number
                                                                                 Tenure
                                                             Referred
                                  Customer
                       Service ID
                                             Count Quarter
                                                                             of
                                                                                      in Offe
                                        ID
                                                              a Friend
                                                                       Referrals
                                                                                Months
                                      8779-
                 IJKDQVSWH3522
                                                         Q3
                                                                              0
                                                 1
                                                                  No
                                                                                       1 None
                                    QRDMV
                                      7495-
                                                                                          Offe
                  BFKMZJAIE2285
             1
                                                 1
                                                         Q3
                                                                  Yes
                                                                              1
                                                                                       8
                                     OOKFY
                                      1658-
                                                                                          Offe
                 EIMVJQBMT7187
                                                 1
                                                         Q3
                                                                  No
                                                                              0
                                                                                      18
                                     BYGOY
                                                                                             Г
                                      4598-
                                                                                          Offe
             3 EROZQXDUU4979
                                                                  Yes
                                                                              1
                                                                                      25
                                                 1
                                                         Q3
                                     XLKNJ
                                      4846-
                                                                                          Offe
                                                                                      37
                  GEEYSJUHY6991
                                                 1
                                                         Q3
                                                                  Yes
                                                                              1
                                    WHAFZ
             •••
                                                          ...
                                                                   ...
                                      2569-
          7038
                 JIVEOZUQQ2296
                                                         Q3
                                                                              0
                                                                                      72
                                                                  No
                                                                                          None
                                    WGERO
                                      6840-
                                                                                          Offe
                                                                                      24
          7039
                 OIQWIUTDY3518
                                                                              1
                                                         Q3
                                                                  Yes
                                     RESVB
                                      2234-
          7040
                 QSHQPZAYF6519
                                                         Q3
                                                                                      72
                                                                  Yes
                                                                                         None
                                    XADUH
                                      4801-
          7041
                 PMJLYZGVQ7211
                                                                              1
                                                         Q3
                                                                  Yes
                                                                                      11
                                                                                          None
                                     JZAZL
                                      3186-
          7042
                   BILLTTPVG8428
                                                         Q3
                                                                  No
                                                                              0
                                                                                      66
                                                                                         None
                                      AJIEK
         7043 rows × 31 columns
In [18]:
           services_df.info()
        <class 'pandas.core.frame.DataFrame'>
        RangeIndex: 7043 entries, 0 to 7042
        Data columns (total 31 columns):
         #
              Column
                                                   Non-Null Count Dtype
         0
              Service ID
                                                   7043 non-null
                                                                    object
         1
              Customer ID
                                                   7043 non-null
                                                                    object
                                                   7043 non-null
         2
              Count
                                                                    int64
         3
              Quarter
                                                   7043 non-null
                                                                    object
              Referred a Friend
                                                   7043 non-null
                                                                    object
              Number of Referrals
                                                   7043 non-null
                                                                    int64
```

```
6
    Tenure in Months
                                       7043 non-null
                                                       int64
    Offer
 7
                                       7043 non-null
                                                       object
 8
    Phone Service
                                       7043 non-null
                                                       object
 9
    Avg Monthly Long Distance Charges 7043 non-null
                                                       float64
10 Multiple Lines
                                                       object
                                       7043 non-null
11 Internet Service
                                       7043 non-null
                                                       object
12 Internet Type
                                       7043 non-null
                                                       object
13 Avg Monthly GB Download
                                       7043 non-null
                                                       int64
14 Online Security
                                       7043 non-null
                                                       object
15 Online Backup
                                       7043 non-null
                                                       object
16 Device Protection Plan
                                       7043 non-null
                                                       object
17 Premium Tech Support
                                       7043 non-null
                                                       object
18 Streaming TV
                                       7043 non-null
                                                       object
19 Streaming Movies
                                       7043 non-null
                                                       object
20 Streaming Music
                                       7043 non-null
                                                       object
21 Unlimited Data
                                       7043 non-null
                                                       object
22 Contract
                                       7043 non-null
                                                       object
23 Paperless Billing
                                       7043 non-null
                                                       object
24 Payment Method
                                       7043 non-null
                                                       object
25 Monthly Charge
                                       7043 non-null
                                                       float64
26 Total Charges
                                       7043 non-null
                                                       float64
27 Total Refunds
                                       7043 non-null
                                                       float64
 28 Total Extra Data Charges
                                       7043 non-null
                                                       int64
 29 Total Long Distance Charges
                                       7043 non-null
                                                       float64
 30 Total Revenue
                                       7043 non-null
                                                       float64
dtypes: float64(6), int64(5), object(20)
```

memory usage: 1.7+ MB

services_df.isna().sum()

In [19]:

```
Out[19]: Service ID
                                                  0
          Customer ID
                                                  0
          Count
                                                  0
                                                  0
          Quarter
          Referred a Friend
                                                  0
          Number of Referrals
                                                  0
          Tenure in Months
                                                  a
          Offer 0 0 1
                                                  0
          Phone Service
                                                  a
          Avg Monthly Long Distance Charges
                                                  a
          Multiple Lines
                                                  0
          Internet Service
                                                  0
          Internet Type
                                                  0
          Avg Monthly GB Download
                                                  0
          Online Security
                                                  0
          Online Backup
                                                  a
```

Device Protection Plan 0 Premium Tech Support 0 Streaming TV 0 Streaming Movies 0 Streaming Music 0 Unlimited Data 0 0 Contract 0 Paperless Billing 0 Payment Method 0 Monthly Charge Total Charges 0 Total Refunds 0 Total Extra Data Charges 0 Total Long Distance Charges 0

Total Revenue dtype: int64

As a final step in our data understanding, we will remove any columns that are unnecessary or contain information that we already have in our other two available dataframes. This way we will prevent overlapping of information and inaccurate feature

0

Δνα

impact on our modeling later. Once again there are many columns in this dataframe that were repeated in the first (or second) dataframe, and therefore are unnecessary. Furthermore, any of the quarter or count information isn't relevant to our modeling process. These things have nothing to do with whether or not a customer churned, because no consumer is keeping track of fiscal quarters of their cellphone company to plan to cut their service, but gender or tenure with company might very well be relevant.

```
In [20]: services_df = services_df.drop(['Service ID', 'Count', 'Quarter', 'Tenure in M
In [21]: services_df
```

Out[21]:

	Customer ID	Referred a Friend	Number of Referrals	Offer	Avg Monthly Long Distance Charges	Internet Type	Avg Monthly GB Download	Streaming Music
0	8779- QRDMV	No	0	None	0.00	DSL	8	No
1	7495- OOKFY	Yes	1	Offer E	48.85	Fiber Optic	17	No
2	1658- BYGOY	No	0	Offer D	11.33	Fiber Optic	52	Yes
3	4598- XLKNJ	Yes	1	Offer C	19.76	Fiber Optic	12	No
4	4846- WHAFZ	Yes	1	Offer C	6.33	Fiber Optic	14	No
•••								
7038	2569- WGERO	No	0	None	22.77	None	0	No
7039	6840- RESVB	Yes	1	Offer C	36.05	Cable	24	Yes
7040	2234- XADUH	Yes	4	None	29.66	Fiber Optic	59	Yes
7041	4801- JZAZL	Yes	1	None	0.00	DSL	17	No
7042	3186- AJIEK	No	0	None	30.96	Fiber Optic	11	Yes

7043 rows × 15 columns

```
In [22]: services_df.info()
```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 15 columns):
Column

Non-Null Count Dtype

In [23]:

```
0
    Customer ID
                                        7043 non-null
                                                        object
1
    Referred a Friend
                                        7043 non-null
                                                        object
2
    Number of Referrals
                                        7043 non-null
                                                        int64
                                                        object
3
                                        7043 non-null
4
   Avg Monthly Long Distance Charges
                                                        float64
                                       7043 non-null
                                       7043 non-null
5
                                                        object
    Internet Type
                                       7043 non-null
6
   Avg Monthly GB Download
                                                        int64
7
   Streaming Music
                                       7043 non-null
                                                        object
8
   Unlimited Data
                                       7043 non-null
                                                        object
9
                                        7043 non-null
                                                        float64
    Monthly Charge
10
                                       7043 non-null
                                                        float64
   Total Charges
11 Total Refunds
                                       7043 non-null
                                                        float64
                                       7043 non-null
                                                        int64
12 Total Extra Data Charges
13 Total Long Distance Charges
                                       7043 non-null
                                                        float64
14 Total Revenue
                                       7043 non-null
                                                        float64
```

dtypes: float64(6), int64(3), object(6)

memory usage: 825.5+ KB

```
services_df.isna().sum()
          Customer ID
                                                0
Out[23]:
          Referred a Friend
                                                0
          Number of Referrals
                                                0
          Offer
                                                0
          Avg Monthly Long Distance Charges
                                                0
          Internet Type
                                                0
          Avg Monthly GB Download
                                                0
          Streaming Music
                                                0
          Unlimited Data
                                                0
```

0 Monthly Charge 0 Total Charges Total Refunds 0 0 Total Extra Data Charges 0 Total Long Distance Charges Total Revenue

dtype: int64

Now, throughout all this data exploration and preparation, we have created three separate dataframes with one common column ("Customer ID"), 7043 entries that correspond to each other in each of the three dataframes, and various information on their services, plans, and habits. At this point, we must weld together those three separate dataframes into one usable one, and then begin the analysis process.

```
In [24]:
          churn_df_2 = churn_df.merge(telco_churn, how="inner", on='Customer ID')
```

In [25]:

churn df 2

Out[25]: Customer Multiple Internet Senior Phone Partner Dependents Tenure Citizen ID Service Lines Service So No 7590-0 No Yes No 1 No phone DSL VHVEG service 5575-1 DSL No No No 34 Yes No **GNVDE** 3668-

No

No

2

Yes

DSL

No

No

7705

QPYBK

No

2

Jackillie	ricii/CiluiiiL	ialloutb	yrıb at mam	men manpyi	CHUITIDACKI			
DSL	phone service	No	45	No	No	No	CFOCW	3
Fiber optic	No	Yes	2	No	No	No	9237- HQITU	4
								•••
DSL	Yes	Yes	24	Yes	Yes	No	6840- RESVB	7038
Fiber optic	Yes	Yes	72	Yes	Yes	No	2234- XADUH	7039
DSL	No phone service	No	11	Yes	Yes	No	4801- JZAZL	7040
Fiber optic	Yes	Yes	4	No	Yes	Yes	8361- LTMKD	7041
Fiber optic	No	Yes	66	No	No	No	3186- AJIEK	7042

7043 rows × 22 columns

In [26]: churn_df_2.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 7043 entries, 0 to 7042
Data columns (total 22 columns):

#	Column	Non-Null Count	Dtype
0	Customer ID	7043 non-null	object
1	Senior Citizen	7043 non-null	object
2	Partner	7043 non-null	object
3	Dependents	7043 non-null	object
4	Tenure	7043 non-null	int64
5	Phone Service	7043 non-null	object
6	Multiple Lines	7043 non-null	object
7	Internet Service	7043 non-null	object
8	Online Security	7043 non-null	object
9	Online Backup	7043 non-null	object
10	Device Protection	7043 non-null	object
11	Tech Support	7043 non-null	object
12	Streaming TV	7043 non-null	object
13	Streaming Movies	7043 non-null	object
14	Contract	7043 non-null	object
15	Paperless Billing	7043 non-null	object
16	Payment Method	7043 non-null	object
17	Churn	7043 non-null	object
18	Gender	7043 non-null	object
19	Churn Value	7043 non-null	int64
20	Churn Score	7043 non-null	int64
21	CLTV	7043 non-null	int64

Our final step in merging these dataframes is to combine the dataframe we just made with the third and final dataframe to create one fully meshed dataset that we can then use for modeling. Because of the steps we took during this preprocessing and

dtypes: int64(4), object(18)
memory usage: 1.2+ MB

exploration, we now have no nulls and properly arranged data, to make the process of modeling much easier going forward.

```
In [27]: final_df = churn_df_2.merge(services_df, how="inner", on='Customer ID')
In [28]: final_df
```

Out[28]:

	Customer ID	Senior Citizen	Partner	Dependents	Tenure	Phone Service	Multiple Lines	Internet Service	Sı
0	7590- VHVEG	No	Yes	No	1	No	No phone service	DSL	
1	5575- GNVDE	No	No	No	34	Yes	No	DSL	
2	3668- QPYBK	No	No	No	2	Yes	No	DSL	
3	7795- CFOCW	No	No	No	45	No	No phone service	DSL	
4	9237- HQITU	No	No	No	2	Yes	No	Fiber optic	
•••									
7038	6840- RESVB	No	Yes	Yes	24	Yes	Yes	DSL	
7039	2234- XADUH	No	Yes	Yes	72	Yes	Yes	Fiber optic	
7040	4801- JZAZL	No	Yes	Yes	11	No	No phone service	DSL	
7041	8361- LTMKD	Yes	Yes	No	4	Yes	Yes	Fiber optic	
7042	3186- AJIEK	No	No	No	66	Yes	No	Fiber optic	

7043 rows × 36 columns

```
In [29]: final_df.info()
```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 7043 entries, 0 to 7042
Data columns (total 36 columns):

#	Column	Non-Null Count	Dtype
0	Customer ID	7043 non-null	object
1	Senior Citizen	7043 non-null	object
2	Partner	7043 non-null	object
3	Dependents	7043 non-null	object
4	Tenure	7043 non-null	int64
_			

```
Phone Service
                                      7043 non-null
                                                     object
  Multiple Lines
                                     7043 non-null
                                                     object
6
   Internet Service
7
                                      7043 non-null
                                                     object
8 Online Security
                                      7043 non-null
                                                     object
   Online Backup
                                      7043 non-null
                                                     object
10 Device Protection
                                      7043 non-null
                                                     object
11 Tech Support
                                      7043 non-null
                                                     object
12 Streaming TV
                                     7043 non-null
                                                     object
13 Streaming Movies
                                     7043 non-null
                                                     object
14 Contract
                                      7043 non-null
                                                     object
15 Paperless Billing
                                      7043 non-null
                                                     object
16 Payment Method
                                      7043 non-null
                                                     object
17 Churn
                                      7043 non-null
                                                     object
18 Gender
                                      7043 non-null
                                                     object
19 Churn Value
                                     7043 non-null
                                                     int64
20 Churn Score
                                      7043 non-null
                                                     int64
21 CLTV
                                      7043 non-null
                                                     int64
22 Referred a Friend
                                      7043 non-null
                                                     object
23 Number of Referrals
                                      7043 non-null
                                                     int64
24 Offer
                                      7043 non-null
                                                     object
25 Avg Monthly Long Distance Charges 7043 non-null
                                                     float64
                                      7043 non-null
26 Internet Type
                                                     object
27 Avg Monthly GB Download
                                     7043 non-null
                                                     int64
28 Streaming Music
                                     7043 non-null
                                                     object
29 Unlimited Data
                                     7043 non-null
                                                     object
                                     7043 non-null
30 Monthly Charge
                                                     float64
31 Total Charges
                                     7043 non-null
                                                     float64
32 Total Refunds
                                     7043 non-null
                                                     float64
33 Total Extra Data Charges
                                      7043 non-null
                                                     int64
34 Total Long Distance Charges
                                      7043 non-null
                                                     float64
35 Total Revenue
                                      7043 non-null
                                                     float64
```

dtypes: float64(6), int64(7), object(23)

memory usage: 2.0+ MB

Et voila! One "perfect" dataframe with many various features that can be used to compare against "Churn" and "Churn Value", respectively, in order to determine how much of an impact those features have on whether a customer will or will not churn. As a final last prepatory step before we continue, we will remove the "Customer ID" column before we begin to model, as it is not necessary and will trip up our model. I have also chosen to remove "CLTV" and "Churn Score" as I was not able to do the extraneous research needed to explain these features and the impact they might have on whether a customer churned or not.

```
In [30]:
          final_df = final_df.drop(['Customer ID', 'Churn Value', 'Churn Score', 'CLTV']
In [31]:
          final df
Out[31]:
```

	Senior Citizen	Partner	Dependents	Tenure	Phone Service	Multiple Lines	Internet Service	Online Security	Or Bac
0	No	Yes	No	1	No	No phone service	DSL	No	
1	No	No	No	34	Yes	No	DSL	Yes	
2	No	No	No	2	Yes	No	DSL	Yes	
3	No	No	No	45	No	No phone	DSL	Yes	

4	No	No	No	2	Yes	No	Fiber optic	No
•••								
7038	No	Yes	Yes	24	Yes	Yes	DSL	Yes
7039	No	Yes	Yes	72	Yes	Yes	Fiber optic	No
7040	No	Yes	Yes	11	No	No phone service	DSL	Yes
7041	Yes	Yes	No	4	Yes	Yes	Fiber optic	No
7042	No	No	No	66	Yes	No	Fiber optic	Yes

7043 rows × 32 columns

```
In [32]: final_df['Churn'].value_counts(normalize=True)
Out[32]: No    0.73463
    Yes    0.26537
    Name: Churn, dtype: float64
```

One very final thing to point out purely to help us better understand our data is that we do have a class imbalance. Of the customers we have available, some 7000 or so, 73% have not churned, while 27% have. This means that as we move into modeling our data, we must account for that class imbalance, because any models with be somewhat "safe" simply predicting that a customer will not churn, and that does not actually help in a business sense at all.

Analysis

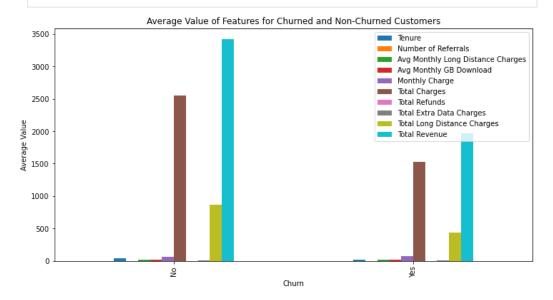
Now it is time to establish correlation between our features (those 35 mentioned) and our target ("Churn") so that we can proceed with modeling those features against that target in order to predict whether a customer will churn with the highest possible recall. Remember that recall is our important metric here rather than accuracy because it measures, in this context, which customers are predicted to churn versus which customers actually churned.

```
in [33]:
    final_df_churn = final_df.groupby("Churn")

# Calculate the average value of each feature, separated by churn status
avg_values = final_df_churn.mean()

# Plot bar plots for each feature
avg_values.plot(kind="bar", figsize=(12, 6))
plt.xlabel("Churn")
plt.ylabel("Average Value")
plt.title("Average Value of Features for Churned and Non-Churned Customers")
plt.legend(loc="upper right")
```

pit.snow()



As can be easily viewed from these bar charts, these features have a varying impact on whether or not a customer will churn, and should provide plenty of usable data for the purpose! From this, we begin to learn which features have the most impact on churn - in this case, total revenue, total charges and total long distance charges seem to be very important factors to whether or not a customer will churn, and that does make a lot of logical sense. If a customer has paid in a lot of money to a company ("Total Revenue"), they would be less likely to churn, as they've invested money. The "sunk cost fallacy" can explain this. To take things a logical step further, customers might also be more likely to churn (or not) if they are consistently charged extra for long distance or other features.

Now that we have established that our features do have some correlation as to whether or not a customer will churn, we can begin the modeling process with what might be our most important step - splitting our data into a test and training set, respectively. We do this because data leakage can cause a model to overperform. In other words, showing the test data to a model whilst it is training means that it is not truly telling us how well it performs on new data.

```
In [34]: X = final_df.drop('Churn', axis=1)
    y = final_df['Churn']
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=.25, rando

In [35]: X_train
```

Out[35]:

	Senior Citizen	Partner	Dependents	Tenure	Phone Service	Multiple Lines		Online Security	
6607	No	No	Yes	1	No	No phone service	DSL	No	
2598	No	No	No	7	Yes	No	Fiber optic	No	

2345	No	No	Yes	4	Yes	No	No	internet service	inte se
4093	No	No	No	29	Yes	Yes	Fiber optic	No	
693	No	No	No	3	Yes	Yes	Fiber optic	No	
•••									
3772	No	Yes	No	1	Yes	No	Fiber optic	Yes	
5191	No	Yes	Yes	23	Yes	Yes	DSL	Yes	
5226	No	Yes	Yes	12	Yes	No	No	No internet service	inte se
5390	Yes	No	No	12	Yes	Yes	Fiber optic	No	
860	No	No	No	26	Yes	No	No	No internet service	inte se

5282 rows × 31 columns

In [36]: X_test

Out[36]:

	Senior Citizen	Partner	Dependents	Tenure	Phone Service	Multiple Lines	Internet Service	Online Security	Or Bac
185	No	Yes	No	1	No	No phone service	DSL	No	
2715	No	No	No	41	Yes	Yes	No	No internet service	int∈ se
3825	No	Yes	Yes	52	Yes	No	No	No internet service	inte se
1807	No	No	No	1	Yes	No	Fiber optic	No	
132	No	No	No	67	Yes	No	DSL	No	
•••									
5845	No	No	No	3	Yes	Yes	Fiber optic	No	
2301	No	Yes	No	8	Yes	Yes	Fiber optic	No	
5121	No	No	No	29	Yes	Yes	DSL	No	
677	Nο	Nο	No	2	Nο	No nhone	DSI	Nο	

service

6062 Yes No No 53 Yes No Fiber No optic

1761 rows × 31 columns

encoded.

Now, we will begin the modeling process by setting up a pipeline in order to scale our numerical features (such as Avg Monthly GB Download) to ensure that they are each taken with the same importance, and encoding our categorical features (such as Internet Service) to allow the model to parse through them properly. Many of our columns are actually Boolean, or "yes/no" data types, but they are listed as objects. Still, they will be able to be used in modeling just fine without any encoding, and we will not require any ordinal encoding as nothing here is in any sort of order. Essentially, the internet type ("Fiber optic", "DSL", or "no") has no listed order; there is no measurable numeric difference between "no" and "Fiber optic", so it does not need to be ordinally

```
In [37]:
           le = LabelEncoder()
In [38]:
           le.fit(y train)
Out[38]:
          LabelEncoder()
In [39]:
           encoded_y_train = le.fit_transform(y_train)
In [40]:
           encoded y = le.transform(y test)
In [41]:
           le.classes
Out[41]: array(['No', 'Yes'], dtype=object)
          In the above cells, I have had to establish a Label Encoder so that our target column,
          "Churn", will be registered as 0/1 rather than "no"/"yes", respectively. This will allow us
          to include the target column in the modeling process, which as I'm sure you can
          imagine is a very important thing to include!
In [42]:
           num_feats = ['Tenure', 'Number of Referrals', 'Avg Monthly Long Distance Charg
           cat_feats = ['Multiple Lines', 'Internet Service', 'Online Security', 'Online
           ord_feats = ['Senior Citizen', 'Partner', 'Dependents', 'Phone Service', 'Pape
          We've encoded "Gender" as male (0) and female (1).
In [43]:
           X train.columns
Out[43]: Index(['Senior Citizen', 'Partner', 'Dependents', 'Tenure', 'Phone Service',
```

'Multiple Lines', 'Internet Service', 'Online Security',

'Dofonnod a Enjoyd'

'Online Backup', 'Device Protection', 'Tech Support', 'Streaming TV', 'Streaming Movies', 'Contract', 'Paperless Billing', 'Payment Method',

'Number of Pofennals'

'Condon'

```
neterreu a rittenu , Number of neterrats ,
                 'Avg Monthly Long Distance Charges', 'Internet Type',
                 'Avg Monthly GB Download', 'Streaming Music', 'Unlimited Data',
                 'Monthly Charge', 'Total Charges', 'Total Refunds',
                 'Total Extra Data Charges', 'Total Long Distance Charges',
                 'Total Revenue'],
                dtype='object')
In [44]:
          prep_pipeline = ColumnTransformer(
              transformers=[
                   ('categorical', OneHotEncoder(handle_unknown='ignore'), cat_feats),
                   ('numeric', StandardScaler(), num_feats),
                  ('other', OrdinalEncoder(categories=[['No', 'Yes'], ['No', 'Yes'], ['N
              ],
              remainder='passthrough')
In [45]:
          X train encoded = prep pipeline.fit transform(X train)
In [46]:
          X test encoded = prep pipeline.transform(X test)
```

What we've done in the above cells is to create a pipeline that will OneHotEncode our categorical features and scale our numeric features, respectively, whilst leaving those that need no attention from the prep pipeline unchanged. At this point, we will instantiate a dummy model and run predictions on our training data. We have also LabelEncoded our target column to transform it from yes/no to 1/0, respectively, in order to make modeling possible.

```
In [47]:
                                                                                                      dummy model = DummyClassifier(strategy="most frequent")
In [48]:
                                                                                                      dummy_model.fit(X_train, y_train)
                                                                                               DummyClassifier(strategy='most_frequent')
 In [49]:
                                                                                                      dummy model.predict(X train)[:50]
Out[49]: array(['No', 'No', 'N
                                                                                                                                                                          'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No'
                                                                                                                                                                       'No', 
                                                                                                                                                                         'No', 'No', 'No', 'No', 'No'], dtype='<U2')
```

As stated, the dummy model has taken the most prevalent input from the 'Churn' target column and made that response its answer each time. In other words, the model has predicted that customers will not churn, because that is the most prevalent value in our target column. Now would be a good time to point out that we do have what is called a class imbalance, meaning that there are far more 'unchurned' customers than 'churned' ones to pull data from. We will be dealing with that throughout the process of modeling and analyzing, but for a dummy model, it is unnecessary to put that work in yet.

```
In [50]:
          y.value counts(normalize=True)
                 0.73463
Out[50]:
          Voc
                 A 26E27
```

```
162 6.50721
```

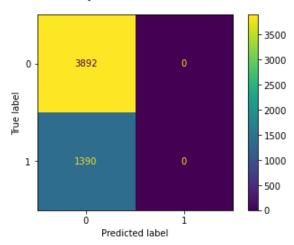
Name: Churn, dtype: float64

With all that in mind, let's see how well our dummy model performed!

```
In [51]: fig, ax = plt.subplots()
    fig.suptitle("Dummy Model Confusion Matrix")
    ConfusionMatrixDisplay(confusion_matrix(y_train, dummy_model.predict(X_train))
```

Out[51]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x2035555122 0>

Dummy Model Confusion Matrix



```
In [52]: cross_val_score(dummy_model, X_train, y_train, cv=5).mean()
```

Out[52]: 0.7368420486797971

At this point, our dummy model has a 73% accuracy score, and it will have 0% recall because it is a dummy model with no variation in its response - it predicts that a customer will not churn and it is correct 73% of the time. Obviously, 73% is not a very good score, and 0% is even worse, because remember that our biggest concern insofar as measurement metrics will be recall for this model. Therefore, the current score to beat is actually 0%, though we would love to see a better accuracy score as well.

First, we will instantiate a simple Decision Tree classifier to see if that alone performs better than our dummy model.

```
In [53]: clf = tree.DecisionTreeClassifier()

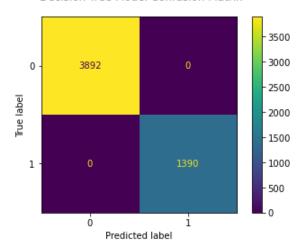
In [54]: clf.fit(X_train_encoded, encoded_y_train)

Out[54]: DecisionTreeClassifier()

In [55]: fig, ax = plt.subplots()
    fig.suptitle("Decision Tree Model Confusion Matrix")
    ConfusionMatrixDisplay(confusion_matrix(encoded_y_train, clf.predict(X_train_encoded_y_train, clf.predict(X_train_encoded_y_train)
```

Out[55]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x20355977a9 0>

Decision Tree Model Confusion Matrix



```
In [56]: cross_val_score(clf, X_train_encoded, encoded_y_train, cv=5).mean()
```

Out[56]: 0.747444884034288

```
In [57]: clf.score(X_train_encoded, encoded_y_train)
```

Out[57]: 1.0

As we can see, our simple decision tree is very overfit. Rather than try to fiddle with the very simple decision tree that will be unlikely to get us anywhere really, I have decided to move into random forest modeling and tweak that instead so that I can find the best model for this job.

```
In [58]: rfc = RandomForestClassifier(random_state=42)
```

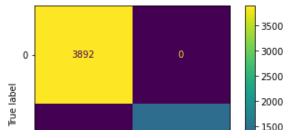
In [59]: rfc.fit(X_train_encoded, encoded_y_train)

Out[59]: RandomForestClassifier(random_state=42)

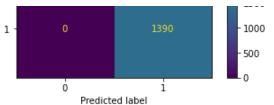
```
fig, ax = plt.subplots()
    fig.suptitle("Random Forest Model Confusion Matrix")
    ConfusionMatrixDisplay(confusion_matrix(encoded_y_train, rfc.predict(X_train_e))
```

Out[60]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x2035544af4 0>

Random Forest Model Confusion Matrix



Out[62]: 1.0



```
In [61]: cross_val_score(rfc, X_train_encoded, encoded_y_train, cv=5).mean()
Out[61]: 0.8152172744474069
In [62]: rfc.score(X_train_encoded, encoded_y_train)
```

As we may have guessed, an unaltered random forest model presents us with the exact same information as an even simpler decision tree model. Now, it's time to start hypertuning our random forest so that we can reduce overfitting and begin to accurately predict whether or not a customer will churn. We will need to tune our model, but also deal with the previously mentioned class imbalance in our data - most of the customers did not churn, and so our models are assuming that most people will not churn.

Another thing to note is that we have done these models outside of a pipeline, mainly because they are very simple models that I expected to both overfit and not predict very well. Thusly, it was not strictly "necessary" to establish a modeling pipeline for these very simple models. We will do that now, and use GridSearchCV to further tune our model.

In the above cells, I have set up some parameter grids to allow us to hypertune our models and choose the best possible parameters to model with in order to achieve high recall and accuracy, though recall is our most important metric.

```
In [66]: grid_search = GridSearchCV(estimator=model_pipeline, param_grid=param_grid, cv
In [67]: grid_search.fit(X_train, encoded_y_train)
```

```
Out[67]: GridSearchCV(cv=5,
                        estimator=Pipeline(steps=[('preprocessing',
                                                      ColumnTransformer(remainder='passthrou
          gh',
                                                                         transformers=[('cate
          gorical',
                                                                                         0neHo
          tEncoder(handle_unknown='ignore'),
                                                                                         ['Mul
          tiple '
                                                                                          'Lin
          es',
                                                                                          'Int
          ernet '
                                                                                          'Ser
          vice',
                                                                                          'Onl
          ine '
                                                                                          'Sec
          urity',
                                                                                          'Onl
          ine '
                                                                                          'Bac
          kup',
                                                                                          'Dev
          ice '
                                                                                          'Pro
          tection',
                                                                                          'Tec
                                                                                          'Sup
          port',
                                                                                          'Str
          eaming '
                                                                                          'Τ
          ۷',
                                                                                          'Str
          eaming '
                                                                                          'Mov
          ies',
                                                                                          'Con
          tract',
                                                                                          'Pay
          ment '
                                                                                          'Met
          hod',
           '0...
                                                                                          'Dep
          endents',
                                                                                          'Pho
          ne '
                                                                                          'Ser
          vice',
                                                                                          'Pap
          erless '
                                                                                          'Bil
          ling',
                                                                                          'Gen
          der',
                                                                                          'Ref
          erred '
                                                                                          'a '
                                                                                          'Fri
          end',
                                                                                          'Str
```

```
eaming '
                                                                                        'Mus
          ic',
                                                                                        'Unl
          imited '
                                                                                        'Dat
          a'])])),
                                                   ('classifier',
                                                    RandomForestClassifier(random state=4
          2))]),
                        param_grid={'classifier__class_weight': [None, 'balanced',
                                                                   'balanced subsample'],
                                    'classifier__max_depth': [None, 5, 10],
                                    'classifier__n_estimators': [100, 200, 300]},
                        scoring='recall')
In [68]:
          fig, ax = plt.subplots()
           fig.suptitle("Random Forest Tuning 1 Confusion Matrix")
           ConfusionMatrixDisplay(confusion_matrix(encoded_y_train, grid_search.predict(X
Out[68]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x20355519fd
          Random Forest Tuning 1 Confusion Matrix
                                                 2500
          0
                  2960
                                                 2000
        Frue label
                                                 - 1500
                                                 1000
                   228
          1
                                                  500
                    ó
                       Predicted label
In [69]:
           grid_search.best_score_
Out[69]: 0.8129496402877698
In [70]:
           grid_search.score(X_train, encoded_y_train)
Out[70]:
          0.8359712230215828
In [71]:
           best_params1 = grid_search.best_params_
           best_model1 = grid_search.best_estimator_
In [72]:
           best params1
Out[72]: {'classifier__class_weight': 'balanced_subsample',
            'classifier__max_depth': 5,
           'classifier__n_estimators': 100}
```

After our first attempt at tuning our model, we have wound up with a slightly lower

accuracy score, but our model is not as overfit. Furthermore, we have achieved a fairly decent recall score. Do recall (ba-dum-tss) that the reason "recall" is our most important metric is that recall shows whether a model can find all objects of the target class - in this case "churn". In other words, can this model find every possible churn case and help us prevent it? Earmarking potential churning customers is far more important than simply accurately predicting that they'll exist.

Now, we must use the hyperparameters given to us by our first attempt at tuning and see if we can't raise our recall, accuracy, or both!

```
In [73]:
           grid_search2 = GridSearchCV(estimator=model_pipeline, param_grid=param_grid2,
In [74]:
           grid_search2.fit(X_train, encoded_y_train)
Out[74]: GridSearchCV(cv=5,
                        estimator=Pipeline(steps=[('preprocessing',
                                                     ColumnTransformer(remainder='passthrou
          gh',
                                                                         transformers=[('cate
          gorical',
                                                                                         OneHo
          tEncoder(handle_unknown='ignore'),
                                                                                         ['Mul
          tiple '
                                                                                          'Lin
          es',
                                                                                          'Int
          ernet '
                                                                                          'Ser
          vice',
                                                                                          'Onl
          ine '
                                                                                          'Sec
          urity',
                                                                                          'Onl
          ine '
                                                                                          'Bac
          kup',
                                                                                          'Dev
          ice '
                                                                                          'Pro
          tection',
                                                                                          'Tec
                                                                                          'Sup
          port',
                                                                                          'Str
          eaming '
                                                                                          'Τ
          ۷',
                                                                                          'Str
          eaming '
                                                                                          'Mov
          ies',
                                                                                          'Con
          tract',
                                                                                          'Pay
          ment '
                                                                                          'Met
          hod',
```

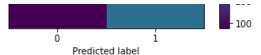
```
'0...
                                                                                           'Par
           tner',
                                                                                           'Dep
           endents',
                                                                                           'Pho
          ne '
                                                                                           'Ser
          vice',
                                                                                           'Pap
           erless '
                                                                                           'Bil
          ling',
                                                                                           'Gen
          der',
                                                                                           'Ref
           erred '
                                                                                           'a '
                                                                                           'Fri
          end',
                                                                                           'Str
          eaming '
                                                                                           'Mus
           ic',
                                                                                           'Unl
           imited '
                                                                                            'Dat
           a'])])),
                                                     ('classifier',
                                                      RandomForestClassifier(random_state=4
           2))]),
                         param_grid={'classifier__class_weight': [None, 'balanced',
                                                                      'balanced_subsample'],
                                      'classifier__max_depth': [4, 5, 6],
                                      'classifier__n_estimators': [90, 100, 110]},
                         scoring='recall')
In [75]:
           fig, ax = plt.subplots()
           fig.suptitle("Random Forest Tuning 2 Confusion Matrix")
           ConfusionMatrixDisplay(confusion_matrix(encoded_y_train, grid_search2.predict(
          <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x203555c4b8</pre>
Out[75]:
           Random Forest Tuning 2 Confusion Matrix
                                                   2500
                   2960
                                                   - 2000
        True label
                                                   - 1500
                                                   1000
                    228
                                                   500
                     ó
                                     i
                        Predicted label
```

hest model? - grid search hest estimator https://github.com/falloutb1tch/ChurnBackTime/blob/main/Final.ipynb

best_params2 = grid_search.best_params_

In [76]:

```
In [77]:
           best params2
          {'classifier__class_weight': 'balanced_subsample',
            classifier__max_depth': 5,
            'classifier n estimators': 100}
In [78]:
           grid search2.best score
Out[78]: 0.8129496402877698
In [79]:
           grid_search2.score(X_train, encoded_y_train)
          0.8359712230215828
Out[79]:
In [80]:
           ypred_final = grid_search2.predict(X_train)
In [81]:
           ypred_final_test = grid_search2.predict(X_test)
In [82]:
           accuracy score(encoded y, ypred final test)
Out[82]: 0.7745599091425327
In [83]:
           recall_score(encoded_y, ypred_final_test)
Out[83]: 0.837160751565762
          Our last model, here referred to as grid_search2, had a recall score of 84 and an
          accuracy score of 77. Due to this, our first tuning attempt turned out to be the more
          successful model, which is referred to as grid_search1. Here is a confusion matrix of
          gridsearch1 on the testing data.
In [98]:
           fig, ax = plt.subplots()
           fig.suptitle("GridSearch1 Test Confusion Matrix")
           ConfusionMatrixDisplay(confusion_matrix(encoded_y, grid_search.predict(X_test)
          <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x20355050f1</pre>
Out[98]:
              GridSearch1 Test Confusion Matrix
                                                   900
                                                   800
                   963
           0
                                                  700
                                                   600
        Frue label
                                                  500
                                                   400
          1
                    78
                                                   300
```



Our last analytic step is to determine feature importance within the model so that we can move forward with advisable next steps for our client. In other words, which of these features impacts customer churn the most? Or the least? Anyway, so I did a bit of a stupid and I did not set up the Capstone environment that we were advised to set up, meaning I am working with an outdated version of sklearn that makes it more difficult to get the feature importances from models. This is my own fault, and also allows me to showcase that I do know how to do it, even the hard way.

```
In [102...
             feature names = final df.columns
In [103...
             feature names
            Index(['Senior Citizen', 'Partner', 'Dependents', 'Tenure', 'Phone Service',
Out[103...
                     'Multiple Lines', 'Internet Service', 'Online Security',
                    'Online Backup', 'Device Protection', 'Tech Support', 'Streaming TV',
                    'Streaming Movies', 'Contract', 'Paperless Billing', 'Payment Method',
                    'Churn', 'Gender', 'Referred a Friend', 'Number of Referrals', 'Offer', 'Avg Monthly Long Distance Charges', 'Internet Type',
                    'Avg Monthly GB Download', 'Streaming Music', 'Unlimited Data',
                    'Monthly Charge', 'Total Charges', 'Total Refunds',
                    'Total Extra Data Charges', 'Total Long Distance Charges',
                    'Total Revenue'],
                   dtype='object')
In [104...
             importance = best model1.named steps['classifier'].feature importances
In [105...
             len(importance)
Out[105...
In [107...
             best model1.named steps['preprocessing']
            ColumnTransformer(remainder='passthrough',
Out[107...
                                transformers=[('categorical',
                                                  OneHotEncoder(handle unknown='ignore'),
                                                  ['Multiple Lines', 'Internet Service',
  'Online Security', 'Online Backup',
                                                   'Device Protection', 'Tech Support',
                                                   'Streaming TV', 'Streaming Movies',
                                                   'Contract', 'Payment Method', 'Offer',
                                                   'Internet Type']),
                                                 ('numeric', StandardScaler(),
                                                  ['Tenure', 'Number of Referrals...
                                                   'Total Extra Data Charges',
                                                   'Total Long Distance Charges',
                                                   'Total Revenue']),
                                                 ('other',
                                                  OrdinalEncoder(categories=[['No', 'Yes'],
                                                                                 ['No', 'Yes'],
['No', 'Yes'],
['No', 'Yes'],
['No', 'Yes'],
                                                                                 ['Male', 'Femal
```

```
e'],
                                                                          ['No', 'Yes'],
['No', 'Yes'],
                                                                          ['No', 'Yes']]),
                                              ['Senior Citizen', 'Partner', 'Dependents',
                                               'Phone Service', 'Paperless Billing',
                                               'Gender', 'Referred a Friend',
                                               'Streaming Music', 'Unlimited Data'])])
In [109...
           ohe_ob = best_model1.named_steps['preprocessing'].named_transformers_['categor'
In [110...
           feature_names = list(ohe_ob.fit(final_df[cat_feats]).get_feature_names())
In [111...
            feature names.extend(num feats)
In [112...
            feature names.extend(ord feats)
In [113...
           len(feature_names)
Out[113...
           60
In [114...
           feature_importance_dict = dict(zip(feature_names, importance))
In [115...
           feature importance dict
Out[115...
           {'x0_No': 0.0007536644088750889,
            'x0_No phone service': 0.00045462930063695333,
            'x0 Yes': 0.0012088887538025575,
            'x1 DSL': 0.010471131528028282,
            'x1_Fiber optic': 0.05605535990829115,
            'x1_No': 0.005828075900586288,
            'x2_No': 0.04308050128351848,
            'x2_No internet service': 0.010758660321975547,
            'x2 Yes': 0.009538193603015714,
            'x3 No': 0.010948717847090852,
            'x3 No internet service': 0.0071142202338824985,
            'x3_Yes': 0.0035219761947465813,
            'x4_No': 0.00550623506937089,
            'x4_No internet service': 0.012410124223649278,
            'x4_Yes': 0.0007705838082981888,
            'x5_No': 0.042529118139103324,
            'x5 No internet service': 0.01793669484824282,
            'x5 Yes': 0.009127110094486211,
            'x6_No': 0.0012623396709702493,
            'x6_No internet service': 0.010379730771566112,
            'x6 Yes': 0.001295866263891002,
            'x7 No': 0.0007621508683868874,
            'x7_No internet service': 0.006295363219339737,
            'x7 Yes': 0.00191331154499985,
            'x8 Month-to-month': 0.12434859554708974,
            'x8 One year': 0.012809908978589822,
            'x8 Two year': 0.08255887793348525,
            100 Bank thancton / automatic \ 1. 0 000000000117210102
```